

Lab 7

Cameron Adams

Math 241, Week 9

```
# Put all necessary libraries here
library(tidyverse)
library(tidytext)
library(gutenbergr)
library(dplyr)
library(tidyr)
library(tidytext)
library(ggplot2)
library(wordcloud)

# Ensure the textdata package is installed
if (!requireNamespace("textdata", quietly = TRUE)) {
  install.packages("textdata")
}
# Load the textdata package
library(textdata)

# Before knitting your document one last time, you will have to download the AFINN lexicon explicitly
lexicon_afinn()
```

```
## # A tibble: 2,477 x 2
##   word      value
##   <chr>    <dbl>
## 1 abandon    -2
## 2 abandoned  -2
## 3 abandons   -2
## 4 abducted   -2
## 5 abduction  -2
## 6 abductions -2
## 7 abhor      -3
## 8 abhorred   -3
## 9 abhorrent  -3
## 10 abhors    -3
## # i 2,467 more rows
```

```
lexicon_nrc()
```

```
## # A tibble: 13,872 x 2
##   word      sentiment
##   <chr>    <chr>
## 1 abacus    trust
## 2 abandon   fear
## 3 abandon   negative
```

```
## 4 abandon      sadness
## 5 abandoned    anger
## 6 abandoned    fear
## 7 abandoned    negative
## 8 abandoned    sadness
## 9 abandonment  anger
## 10 abandonment fear
## # i 13,862 more rows
```

Due: Friday, March 29th at 5:30pm

Goals of this lab

1. Practice matching patterns with regular expressions.
2. Practice manipulating strings with `stringr`.
3. Practice tokenizing text with `tidytext`.
4. Practice looking at word frequencies.
5. Practice conducting sentiment analysis.

Problem 1: What's in a Name? (You'd Be Surprised!)

1. Load the `babynames` dataset, which contains yearly information on the frequency of baby names by sex and is provided by the US Social Security Administration. It includes all names with at least 5 uses per year per sex. In this problem, we are going to practice pattern matching!

```
library(babynames)
data("babynames")
#?babynames
```

- a. For 2000, find the ten most popular female baby names that start with the letter Z.

```
female_names_2000_Z <- babynames %>%
  filter(year == 2000, sex == "F", startsWith(name, "Z"))

top_female_names_2000_Z <- female_names_2000_Z %>%
  arrange(desc(n)) %>%
  top_n(10)

top_female_names_2000_Z
```

- b. For 2000, find the ten most popular female baby names that contain the letter z.

```
female_names_2000_C_Z <- babynames %>%
  filter(year == 2000, sex == "F", str_detect(name, "z") | startsWith(name, "Z") | str_sub(name, -1) == "z")

top_female_names_2000_C_Z <- female_names_2000_C_Z %>%
  arrange(desc(n)) %>%
  top_n(10)

top_female_names_2000_C_Z
```

```
## # A tibble: 10 x 5
##   year sex  name      n    prop
##   <dbl> <chr> <chr>   <int>  <dbl>
## 1  2000 F    Elizabeth 15094 0.00757
## 2  2000 F    Mackenzie  6348 0.00318
## 3  2000 F     Zoe       3785 0.00190
```

```
## 4 2000 F Mckenzie 2526 0.00127
## 5 2000 F Makenzie 1613 0.000809
## 6 2000 F Jazmin 1391 0.000697
## 7 2000 F Jazmine 1353 0.000678
## 8 2000 F Lizbeth 817 0.000410
## 9 2000 F Eliza 759 0.000380
## 10 2000 F Litzy 722 0.000362
```

c. For 2000, find the ten most popular female baby names that end in the letter z.

```
female_names_2000_end_Z <- babynames %>%
  filter(year == 2000, sex == "F", str_sub(name, -1) == "z")

top_female_names_2000_end_Z <- female_names_2000_end_Z %>%
  arrange(desc(n)) %>%
  top_n(10)

top_female_names_2000_end_Z
```

```
## # A tibble: 11 x 5
##   year sex name n prop
##   <dbl> <chr> <chr> <int> <dbl>
## 1 2000 F Luz 489 0.000245
## 2 2000 F Beatriz 357 0.000179
## 3 2000 F Mercedes 141 0.0000707
## 4 2000 F Maricruz 96 0.0000481
## 5 2000 F Liz 72 0.0000361
## 6 2000 F Inez 69 0.0000346
## 7 2000 F Odaliz 24 0.0000120
## 8 2000 F Marycruz 23 0.0000115
## 9 2000 F Cruz 19 0.00000952
## 10 2000 F Deniz 16 0.00000802
## 11 2000 F Taiz 16 0.00000802
```

d. Between your three tables in 1.a - 1.c, do any of the names show up on more than one list? If so, which ones? (Yes, I know you could do this visually but use some joins!)

```
joined_tables1 <- inner_join(top_female_names_2000_C_Z,
  , top_female_names_2000_C_Z, by = "name")

joined_tables2 <- inner_join(top_female_names_2000_C_Z,
  , top_female_names_2000_end_Z, by = "name")

joined_tables3 <- inner_join(top_female_names_2000_C_Z,
  , top_female_names_2000_C_Z, by = "name")

common_names1 <- joined_tables1$name
common_names2 <- joined_tables2$name
common_names3 <- joined_tables3$name

common_names1
```

```
## [1] "Elizabeth" "Mackenzie" "Zoe" "Mckenzie" "Makenzie" "Jazmin"
## [7] "Jazmine" "Lizbeth" "Eliza" "Litzy"
```

```
common_names2
```

```
## character(0)
```

```
common_names3
```

```
## [1] "Elizabeth" "Mackenzie" "Zoe"          "Mckenzie"  "Makenzie"  "Jazmin"
## [7] "Jazmine"   "Lizbeth"   "Eliza"     "Litzy"
```

From here we can see Zoe is the only one that is shared between lists.

e. Verify that none of the baby names contain a numeric (0-9) in them.

```
names_with_numeric <- babynames %>%
  filter(str_detect(name, "[0-9]")) %>%
  pull(name)

if (length(names_with_numeric) > 0) {
  cat("Baby names containing numeric digits:\n")
  cat(names_with_numeric, sep = "\n")
} else {
  print("No baby names contain numeric digits.")
}
```

```
## [1] "No baby names contain numeric digits."
```

f. While none of the names contain 0-9, that doesn't mean they don't contain "one", "two", ..., or "nine". Create a table that provides the number of times a baby's name contained the word "zero", the word "one", ... the word "nine".

Notes:

- I recommend first converting all the names to lower case.
- If none of the baby's names contain the written number, there you can leave the number out of the table.
- Use `str_extract()`, not `str_extract_all()`. (We will ignore names where more than one of the words exists.)

Hint: You will have two steps that require pattern matching: 1. Subset your table to only include the rows with the desired words. 2. Add a column that contains the desired word.

```
#lowercase
babynames_lower <- babynames %>%
  mutate(name = tolower(name))

numbers_words <- c("zero", "one", "two", "three", "four", "five", "six", "seven", "eight", "nine")

extracted_numbers <- sapply(numbers_words, function(word) {
  count <- babynames_lower %>%
    filter(str_detect(name, paste0("\\b", word, "\\b"))) %>%
    summarise(count = n())
  count <- ifelse(is.na(count$count), 0, count$count)
  return(count)
})

numbers_table <- data.frame(Number = numbers_words, Count = extracted_numbers)

print(numbers_table)
```

```
##      Number Count
## zero    zero    1
## one     one     0
## two     two     0
## three   three   0
## four    four    0
## five    five    0
## six     six     0
## seven   seven   49
## eight   eight   0
## nine    nine    0
```

g. Which written number or numbers don't show up in any of the baby names?

the numbers zero and seven showed up, with 7 showing up 49 times and zero showing up once

h. Create a table that contains the names and their frequencies for the two least common written numbers.

```
babynames_lower <- babynames %>%
  mutate(name = tolower(name))

least_common_numbers <- c("zero", "seven")

least_common_names <- list()

for (number in least_common_numbers) {
  # Filter the dataset for names containing the current number
  filtered_data <- babynames_lower %>%
    filter(str_detect(name, paste0("\\b", number, "\\b")))

  summarized_data <- filtered_data %>%
    group_by(name) %>%
    summarize(frequency = n()) %>%
    arrange(frequency)

  top_names <- head(summarized_data, 2)

  least_common_names[[number]] <- top_names
}

least_common_table <- do.call(rbind, least_common_names)

print(least_common_table)
```

```
## # A tibble: 2 x 2
##   name frequency
## * <chr>      <int>
## 1 zero         1
## 2 seven       49
```

i. List out the names that contain no vowels (consider "y" to be a vowel).

```
library(stringr)

remove_vowels <- function(name) {
  str_remove_all(tolower(name), "[aeiouy]")
}
```

```

babynames_processed <- babynames %>%
  mutate(processed_name = remove_vowels(name))

names_grouped_count <- babynames_processed %>%
  group_by(processed_name) %>%
  summarise(count = n()) %>%
  filter(processed_name != "") # Filter out empty processed names

print(names_grouped_count)

## # A tibble: 17,943 x 2
##   processed_name count
##   <chr>          <int>
## 1 b             1463
## 2 bb            1618
## 3 bbb           811
## 4 bbbj          118
## 5 bbbjn           7
## 6 bbbbl           5
## 7 bbbblnn         1
## 8 bbbbls          1
## 9 bbbr           20
## 10 bbbs           6
## # i 17,933 more rows

```

Problem 2: Tidying the “Call of the Wild”

Did you read “Call of the Wild” by Jack London? If not, [read the first paragraph of its wiki page](#) for a quick summary and then let’s do some text analysis on this classic! The following code will pull the book into R using the `gutenbergr` package.

```

library(gutenbergr)
wild <- gutenberg_download(215)

```

- a. Create a tidy text dataset where you tokenize by words.

```

wild_tokens <- wild %>%
  unnest_tokens(word, text)

stop_words <- get_stopwords("en")
wild_tokens <- wild_tokens %>%
  anti_join(stop_words)

```

- b. Find the frequency of the 20 most common words. First, remove stop words.

```

top_words <- wild_tokens %>%
  count(word, sort = TRUE) %>%
  slice(1:20)

head(top_words)

```

```

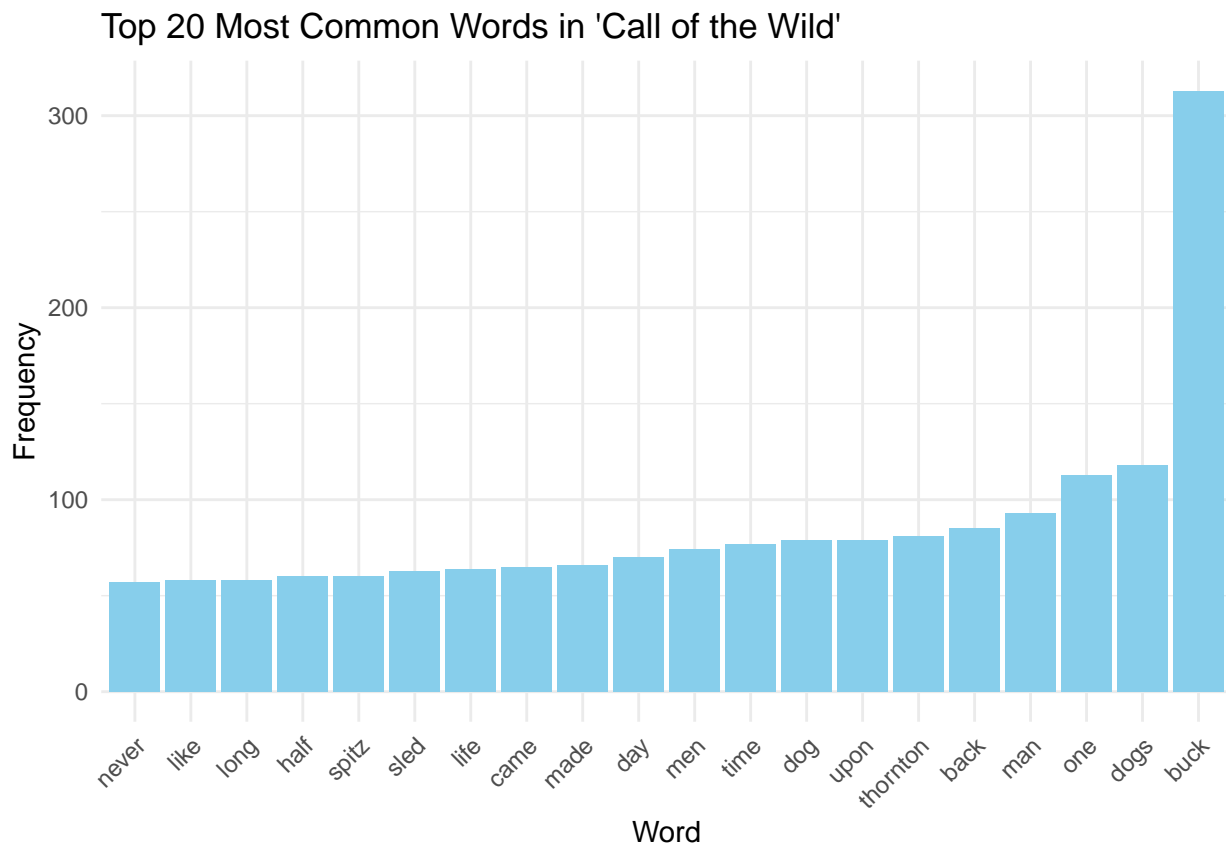
## # A tibble: 6 x 2
##   word      n
##   <chr>    <int>
## 1 buck     313
## 2 dogs     118

```

```
## 3 one      113
## 4 man      93
## 5 back     85
## 6 thornton 81
```

c. Create a bar graph and a word cloud of the frequencies of the 20 most common words.

```
ggplot(top_words, aes(x = reorder(word, n), y = n)) +
  geom_bar(stat = "identity", fill = "skyblue") +
  labs(title = "Top 20 Most Common Words in 'Call of the Wild'",
       x = "Word", y = "Frequency") +
  theme_minimal() +
  theme(axis.text.x = element_text(angle = 45, hjust = 1))
```



```
library(RColorBrewer)
pal <- brewer.pal(9, "Set1")

wordcloud(words = top_words$word, freq = top_words$n, scale=c(3,0.5), min.freq=1,
          max.words=Inf, random.order=FALSE, rot.per=0.35, colors=brewer.pal(8, "Dark2"))
```



- d. Explore the sentiment of the text using three of the sentiment lexicons in `tidytext`. What does your analysis say about the sentiment of the text?

Notes:

- Make sure to NOT remove stop words this time.
- `afinn` is a numeric score and should be handled differently than the categorical scores.

```
wild_tokens <- wild %>%
  unnest_tokens(word, text)

wild_sentiment_afinn <- wild_tokens %>%
  inner_join(get_sentiments("afinn"), by = "word")

wild_sentiment_bing <- wild_tokens %>%
  inner_join(get_sentiments("bing"), by = "word")

wild_sentiment_nrc <- wild_tokens %>%
  inner_join(get_sentiments("nrc"), by = "word")

wild_sentiment_bing <- wild_sentiment_bing %>%
  mutate(score = ifelse(sentiment == "positive", 1, ifelse(sentiment == "negative", -1, 0)))

summary_bing <- summarise(wild_sentiment_bing, sentiment = sum(score))

print("Summary of Bing Sentiment:")

## [1] "Summary of Bing Sentiment:"
print(summary_bing)

## # A tibble: 1 x 1
##   sentiment
##   <dbl>
## 1      -675

print("Summary of Bing Sentiment:")

## [1] "Summary of Bing Sentiment:"
print(summary_bing)

## # A tibble: 1 x 1
##   sentiment
##   <dbl>
## 1      -675
```



```
print("Summary of NRC Sentiment:")
```

```
## [1] "Summary of NRC Sentiment:"
```

```
head(wild_sentiment_nrc)
```

```
## # A tibble: 6 x 3
```

```
##   gutenbergs_id word      sentiment
##   <int> <chr>      <chr>
## 1      215 cover      trust
## 2      215 wild       negative
## 3      215 wild       surprise
## 4      215 law        trust
## 5      215 fang       fear
## 6      215 dominant   fear
```

- e. If you didn't do so in 2.d, compute the average sentiment score of the text using `afinn`. Which positive words had the biggest impact? Which negative words had the biggest impact?

```
wild_sentiment_afinn_filtered <- wild_sentiment_afinn %>%
  filter(!is.na(value))
```

```
avg_sentiment_afinn <- mean(wild_sentiment_afinn_filtered$value)
```

```
print(paste("Average sentiment score (AFINN):", avg_sentiment_afinn))
```

```
## [1] "Average sentiment score (AFINN): -0.322246455834242"
```

```
top_positive_words <- head(wild_sentiment_afinn_filtered[order(wild_sentiment_afinn_filtered$value, decreasing = TRUE)], 5)
top_negative_words <- head(wild_sentiment_afinn_filtered[order(wild_sentiment_afinn_filtered$value)], 5)
```

```
print("Top positive words:")
```

```
## [1] "Top positive words:"
```

```
print(top_positive_words$word)
```

```
## [1] "miracle" "miracle" "wonderful" "triumph" "win"
```

```
print("Top negative words:")
```

```
## [1] "Top negative words:"
```

```
print(top_negative_words$word)
```

```
## [1] "cock" "hell" "hell" "hell" "hell"
```

- f. You should have found that “no” was an important negative word in the sentiment score. To know if that really makes sense, let's turn to the raw lines of text for context. Pull out all of the lines that have the word “no” in them. Make sure to not pull out extraneous lines (e.g., a line with the word “now”).

```
lines_with_no <- wild %>%
  filter(grepl("\\bno\\b", text, ignore.case = TRUE))
```

```
print(lines_with_no$text)
```

```
## [1] "Manuel's treachery. No one saw him and Buck go off through the orchard"
## [2] "solitary man, no one saw them arrive at the little flag station known"
## [3] "that it was the club, but his madness knew no caution. A dozen times he"
```

[4] "'He's no slouch at dog-breakin', that's wot I say," one of the men on"
[5] "all, that he stood no chance against a man with a club. He had learned"
[6] "in the red sweater. "And seem' it's government money, you ain't got no"
[7] "animal. The Canadian Government would be no loser, nor would its"
[8] "while he developed no affection for them, he none the less grew"
[9] "The other dog made no advances, nor received any; also, he did not"
[10] "heart of civilization and flung into the heart of things primordial. No"
[11] "knew no law but the law of club and fang."
[12] "full-grown wolf, though not half so large as she. There was no warning,"
[13] "Buck to trouble him in his sleep. So that was the way. No fair play."
[14] "tail appeasingly, turned to run when he saw that appeasement was of no"
[15] "his flank. But no matter how Spitz circled, Joe whirled around on his"
[16] "their comradeship had no more trouble. His only apparent ambition, like"
[17] "again he returned. Were they in the tent? No, that could not be, else"
[18] "unduly civilized dog, and of his own experience knew no trap and so"
[19] "no ice at all."
[20] "him of his unfinished ration. There was no defending it. While he was"
[21] "days, no matter what the odds, he had never run from a fight. But the"
[22] "internal as well as external economy. He could eat anything, no matter"
[23] "wind and forecast it a night in advance. No matter how breathless the"
[24] "they ran it down. It was no task for him to learn to fight with cut and"
[25] "he betrayed no impatience, shunned all offensive acts."
[26] "as he circled back and forth for a chance to spring in. Buck was no"
[27] "less eager, and no less cautious, as he likewise circled back and forth"
[28] "terrifying, irresistible. There was no opposing them. The team-dogs"
[29] "huskies, there was no hope for him. But he braced himself to the shock"
[30] "nothing, no matter how remotely eatable, had escaped them. They had"
[31] "Again, the rim ice broke away before and behind, and there was no"
[32] "Things no longer went right. There was continual bickering and"
[33] "had destroyed the solidarity of the team. It no longer was as one dog"
[34] "into all kinds of petty misdemeanors. No more was Spitz a leader"
[35] "There was no hope for him. Buck was inexorable. Mercy was a thing"
[36] "make good time. No more Spitz, no more trouble, sure."
[37] "trail. There was no place for Buck save at the front. Once more"
[38] "'Nevaire such a dog as dat Buck!" he cried. "No, nevaire! Heem worth"
[39] "there was no new-fallen snow with which to contend. It was not too"
[40] "Dawson. It was no light running now, nor record time, but heavy toil"
[41] "dim and distant, and such memories had no power over him. Far more"
[42] "ate, and no man sought his sleeping-robe till he had seen to the feet"
[43] "but they could locate no broken bones, could not make it out."
[44] "The half-breed tried to drive him away with the whip; but he paid no"
[45] "They were all terribly footsore. No spring or rebound was left in them."
[46] "slow and prolonged strength drainage of months of toil. There was no"
[47] "power of recuperation left, no reserve strength to call upon. It had"
[48] "manner, but no businesslike method. The tent was rolled into an awkward"
[49] "lead. Hal cried "Whoa! whoa!" but they gave no heed. He tripped and was"
[50] "nothing lively about it, no snap or go in him and his fellows. They"
[51] "Buck felt vaguely that there was no depending upon these two men and"
[52] "at all. And on no day did they succeed in making more than half the"
[53] "that for love or money no additional dog-food was to be obtained. So he"
[54] "the ration of the husky, so the six Outside dogs under Buck could do no"
[55] "had no inkling of such a patience. They were stiff and in pain; their"
[56] "sex-prerogative, she made their lives unendurable. She no longer"
[57] "nightmare. He pulled when he could; when he could no longer pull, he"

```

## [58] "fresher; and Buck, still at the head of the team, but no longer"
## [59] "response to Thornton's warning to take no more chances on the rotten"
## [60] "made no effort. He lay quietly where he had fallen. The lash bit into"
## [61] "pain left him. He no longer felt anything, though very faintly he could"
## [62] "hear the impact of the club upon his body. But it was no longer his"
## [63] "Thornton stood between him and Buck, and evinced no intention of"
## [64] "Hal had no fight left in him. Besides, his hands were full with his"
## [65] "To Buck's surprise these dogs manifested no jealousy toward him. They"
## [66] "names. Buck knew no greater joy than that rough embrace and the sound"
## [67] "had come into the Northland had bred in him a fear that no master could"
## [68] "but the strange dog, no matter what the breed or valor, swiftly"
## [69] "knew there was no middle course. He must master or be mastered; while"
## [70] "Thornton shook his head. "No, it is splendid, and it is terrible, too."
## [71] "the rapids, a stretch of wild water in which no swimmer could live."
## [72] "permitting no slack, while Pete kept it clear of coils. Buck held on"
## [73] "no thousand dollars; nor had Hans or Pete."
## [74] "There were no takers. Not a man believed him capable of the feat."
## [75] "heavy fore legs were no more than in proportion with the rest of the"
## [76] "frankly down his cheeks. "Sir," he said to the Skookum Bench king, "no,"
## [77] "in tragedy and shrouded in mystery. No one knew of the first man. The"
## [78] "But no living man had looted this treasure house, and the dead were"
## [79] "Being in no haste, Indian fashion, he hunted his dinner in the course"
## [80] "uncharted vastness, where no men were and yet where men had been if the"
## [81] "wildfowl had been, but where then there was no life nor sign of"
## [82] "packed flat, And that was all-no hint as to the man who in an early day"
## [83] "washing-pan. They sought no farther. Each day they worked earned them"
## [84] "He had made no noise, yet it ceased from its howling and tried to sense"
## [85] "that no harm was intended, finally sniffed noses with him. Then they"
## [86] "again he took to wandering in the woods, but the wild brother came no"
## [87] "behind who would quarrel no more."
## [88] "secrecy of the forest. He no longer marched. At once he became a thing"
## [89] "proceeded to cut the bull out from the herd. It was no slight task. He"
## [90] "been there throughout the summer. No longer was this fact borne in upon"
## [91] "next bound tearing wide the throat of a second man. There was no"
## [92] "from which no trace led away."
## [93] "was harder to kill a husky dog than them. They were no match at all,"
## [94] "claims of man no longer bound him."

```

g. Draw some conclusions about how “no” is used in the text.

for almost all cases no is used as a negation, not as a negative word. So for example no longer, no slight task, no longer marched, quarrel no more, no men , no slack, no thousand dollars. These are all used to negate the discription.

h. We can also look at how the sentiment of the text changes as the text progresses. Below, I have added two columns to the original dataset. Now I want you to do the following wrangling:

- Tidy the data (but don't drop stop words).
- Add the word sentiments using `bing`.
- Count the frequency of sentiments by index.
- Reshape the data to be wide with the count of the negative sentiments in one column and the positive in another, along with a column for index.
- Compute a sentiment column by subtracting the negative score from the positive.

```

wild_time <- wild %>%
  mutate(line = row_number(), index = floor(line/45) + 1)

```


```
#Hint: fill = 0 will insert zero instead of NA  
#pivot_XXX(..., values_fill = 0)
```

```
wild_tokens <- wild_time %>%  
  unnest_tokens(word, text)  
  
wild_sentiment_bing <- wild_tokens %>%  
  inner_join(get_sentiments("bing"), by = "word")  
  
sentiment_counts <- wild_sentiment_bing %>%  
  count(index, sentiment)  
  
sentiment_counts_wide <- sentiment_counts %>%  
  pivot_wider(names_from = sentiment, values_from = n, values_fill = 0)  
  
sentiment_counts_wide <- sentiment_counts_wide %>%  
  mutate(sentiment = positive - negative)  
  
print(sentiment_counts_wide)
```

```
## # A tibble: 69 x 4  
##   index negative positive sentiment  
##   <dbl>     <int>     <int>     <int>  
## 1     1         9         7        -2  
## 2     2         9        15         6  
## 3     3        26        17        -9  
## 4     4        17         6       -11  
## 5     5        29        10       -19  
## 6     6        25         8       -17  
## 7     7        22        14        -8  
## 8     8        12        11         -1  
## 9     9        20        19         -1  
## 10    10        24        11       -13  
## # i 59 more rows
```

i. Create a plot of the sentiment scores as the text progresses.

```
ggplot(sentiment_counts_wide, aes(x = index, y = sentiment)) +  
  geom_line() +  
  labs(title = "Sentiment Scores as Text Progresses",  
        x = "Index",  
        y = "Sentiment Score")
```



- j. The choice of 45 lines per chunk was pretty arbitrary. Try modifying the index value a few times and recreating the plot in i. Based on your plots, what can you conclude about the sentiment of the novel as it progresses?

```
wild_time_modified <- wild %>%
  mutate(line = row_number(), index = floor(line/30) + 1)

wild_tokens_modified <- wild_time_modified %>%
  unnest_tokens(word, text)

wild_sentiment_bing_modified <- wild_tokens_modified %>%
  inner_join(get_sentiments("bing"), by = "word")


sentiment_counts_modified <- wild_sentiment_bing_modified %>%
  count(index, sentiment)

sentiment_counts_wide_modified <- sentiment_counts_modified %>%
  pivot_wider(names_from = sentiment, values_from = n, values_fill = 0)

sentiment_counts_wide_modified <- sentiment_counts_wide_modified %>%
  mutate(sentiment = positive - negative)

ggplot(sentiment_counts_wide_modified, aes(x = index, y = sentiment)) +
  geom_line() +
  labs(title = "Sentiment Scores as Text Progresses (Modified Index)",
       x = "Index",
       y = "Sentiment Score")
```

Sentiment Scores as Text Progresses (Modified Index)



k. Let's look at the bigrams (2 consecutive words). Tokenize the text by bigrams.

```
wild_bigrams <- wild %>%  
  unnest_tokens(bigram, text, token = "ngrams", n = 2)  
  
head(wild_bigrams)
```

```
## # A tibble: 6 x 2  
##   gutenber_id bigram  
##   <int> <chr>  
## 1      215 <NA>  
## 2      215 <NA>  
## 3      215 <NA>  
## 4      215 <NA>  
## 5      215 <NA>  
## 6      215 the call
```

l. Produce a sorted table that counts the frequency of each bigram and notice that stop words are still an issue.

```
# Load required libraries  
library(dplyr)  
library(tidytext)  
  
wild_bigrams <- wild %>%  
  unnest_tokens(bigram, text, token = "ngrams", n = 2)
```

```

stop_words <- get_stopwords("en")
wild_bigrams <- wild_bigrams %>%
  anti_join(stop_words, by = c("bigram" = "word"))

bigram_counts <- wild_bigrams %>%
  count(bigram, sort = TRUE)

print("Sorted table of bigram frequencies:")

## [1] "Sorted table of bigram frequencies:"
print(bigram_counts)

```

```

## # A tibble: 18,908 x 2
##   bigram      n
##   <chr>   <int>
## 1 <NA>     408
## 2 of the   233
## 3 in the   172
## 4 he was   127
## 5 to the   116
## 6 it was   107
## 7 and the    95
## 8 on the    80
## 9 he had    77
## 10 at the   68
## # i 18,898 more rows

```